

D.1: Model Interpretability

Dr. Bean - Stat 5100

In what ways is a linear model “interpretable”? In what ways is it NOT interpretable?

Why should we care about interpretability?

How might a statistical model be racist/prejudiced? Are linear models also prone to prejudice? If so, how?

How might statistics be used to ensure that the predictive models of tomorrow don't suffer from the prejudices of yesterday?