4.1.1: SAS - Penalized Regression Methods
(Ridge Regression, LASSO, and Elastic Net)

Example 1: (Ridge Regression; recall Handout 2.6.1 example) A study seeks to relate (in females) amount of body fat (Y) to triceps skinfold thickness (X_1), thigh circumference (X_2), and midarm circumference (X_3). Amount of body fat is expensive to measure, requiring immersion of person in water. This expense motivates the desire for a predictive model based on these inexpensive predictors.

```
/* Input data */
data bodyfat;
   input triceps thigh midarm body @@; cards;
  19.5  43.1  29.1  11.9      24.7  49.8  28.2  22.8
  30.7  51.9  37.0  18.7      29.8  54.3  31.1  20.1
  19.1  42.2  30.9  12.9      25.6  53.9  23.7  21.7
  31.4  58.5  27.6  27.1      27.9  52.1  30.6  25.4
  22.1  49.9  23.2  21.3      25.5  53.5  24.8  19.3
  31.1  56.6  30.0  25.4      30.4  56.7  28.3  27.2
  18.7  46.5  23.0  11.7      19.7  44.2  28.6  17.8
  14.6  42.7  21.3  12.8      29.5  54.4  30.1  23.9
  27.7  55.3  25.7  22.6      30.2  58.6  24.6  25.4
  22.7  48.2  27.1  14.8      25.2  51.0  27.5  21.1
;
run;

/* Look at original fit */
proc reg data=bodyfat;
  model body = triceps thigh midarm / vif;
  title1 'Bodyfat Regression (original fit)';
run;
```

Bodyfat Regression (original fit)

Parameter Estimates

Variable	DF	Parameter Estimate	Standard Error	t Value	Pr > |t|	Variance Inflation
Intercept	**1**	117.08469	99.78240	1.17	0.2578	0
triceps	**1**	4.33409	3.01551	1.44	0.1699	708.84291
thigh	**1**	-2.85685	2.58202	-1.11	0.2849	564.34339
midarm	**1**	-2.18606	1.59550	-1.37	0.1896	104.60601

```
/* Try ridge regression as a remedial measure */
proc reg data=bodyfat ridge=0 to .04 by .005
     outvif outest=ridgests
     plots(only)=ridge(VIFaxis=log);
  model body = triceps thigh midarm / vif;
  title1 'Bodyfat Ridge Regression';
  title2 '(find adequate ridge parameter)';
run;
/* What these options do:

     ridge=0 to .04 by 0.005
        run a regression with each of these ridge parameter
       values

     outvif outest=ridgests
        ask for relevant output to be sent to a data set
       called ridgests (will include VIF and standardized
       coefficients for each ridge parameter)

     plots(only)=ridge(VIFaxis=log);
       make Ridge Trace and VIF plots only, with vertical axis
       in VIF plot on log scale
 */
```

Bodyfat Ridge Regression
(find adequate ridge parameter)

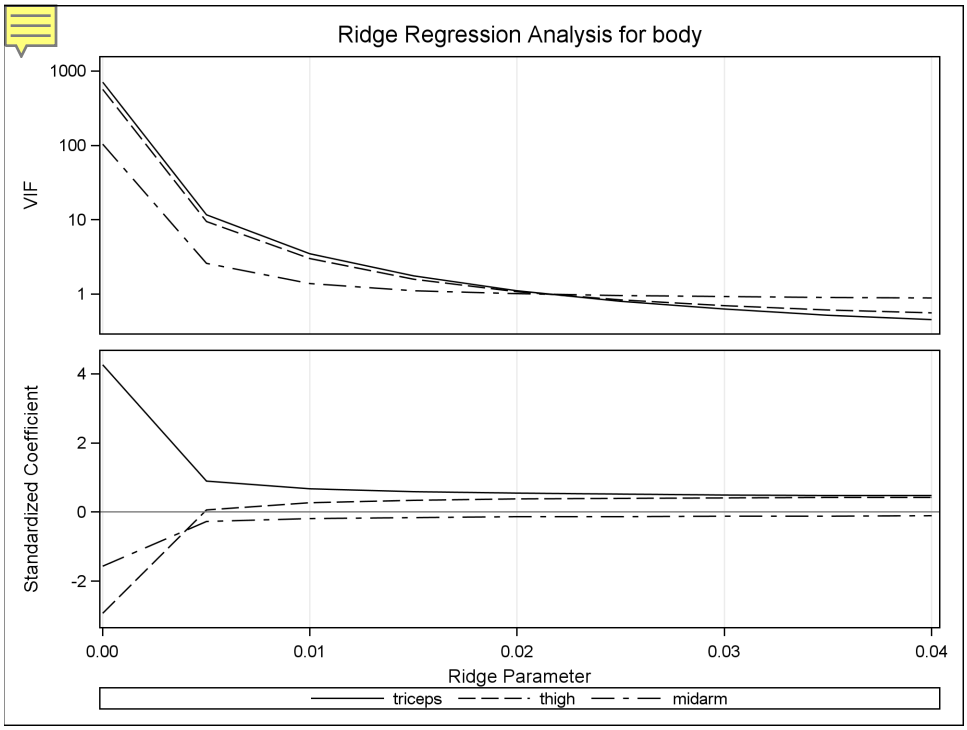

```
/* Now look at variable coeffs with ridge parameter 0.02 */
proc reg data=bodyfat outest=ridgenew outseb ridge=0.02
       outvif noprint;
    model body = triceps thigh midarm;
    title1 'Bodyfat Ridge Regression (c=.02)';
run;
proc print data=ridgenew;
  var _type_  _rmse_  triceps thigh midarm;
  title1 'Ridge Estimates for Variable Coefficients,';
  title2 'with ridge parameter c = 0.02';
run;
/* PARMS and SEB give the result of the regular OLS regression.
   RIDGE and RIDGESEB give the result of the ridge regression.
   -- Note no intercept is given; need to use textbook
      equation 7.46b to get intercept in ridge reg. (as below)
   Note substantial drop in SE for estimates in ridge reg.
   RIDGEVIF give the VIF after ridge regression.
 */
```

Ridge Estimates for Variable Coefficients,
with ridge parameter c = 0.02

Obs	_TYPE_	_RMSE_	triceps	thigh	midarm
1	PARMS	2.47998	4.33409	-2.85685	-2.18606
2	SEB	2.47998	3.01551	2.58202	1.59550
3	RIDGEVIF	.	1.10255	1.08054	1.01051
4	RIDGE	2.59924	0.55535	0.36814	-0.19163
5	RIDGESEB	2.59924	0.12465	0.11841	0.16436

```
/* Get intercept term in ridge regression */
proc means data=bodyfat mean;
  var body triceps thigh midarm;
  title1 'Summary Statistics';
run;
data temp;
 b0 = 20.195 - 0.55535*25.305 - 0.36814*51.17 + 0.19163*27.62;
proc print data=temp;
 var b0;
 title1 'Ridge Regression Intercept';
run;
```

Summary Statistics

Variable	Mean
body	20.1950000
triceps	25.3050000
thigh	51.1700000
midarm	27.6200000

Ridge Regression Intercept

Obs	b0
1	-7.40303

```
/* Get predicted values in ridge regression */
data bodyfat; set bodyfat;
  predbody = -7.40303 + 0.55535*triceps
             + 0.36814*thigh - 0.19163*midarm;
proc sgplot data=bodyfat;
  scatter x=body y=predbody / markerattrs=(symbol=CIRCLEFILLED);
  xaxis label='Observed Body Fat' labelattrs=(size=15pt);
  yaxis label='Predicted Body Fat' labelattrs=(size=15pt);
  title1;
run;
```

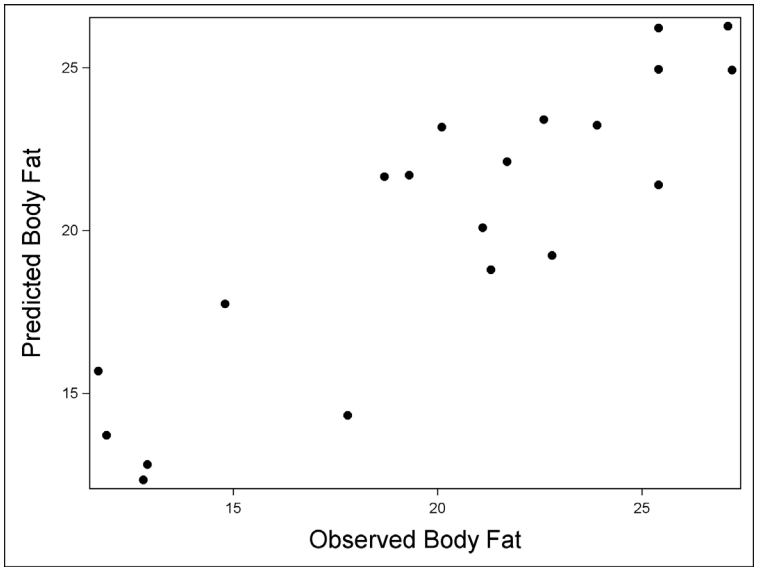

Example 2: (Baseball) This data set (from the SAS Help) contains salary (for 1987) and performance (1986 and some career) data for 322 MLB players who played at least one game in both 1986 and 1987 seasons, excluding pitchers. How can salary be predicted from performance?

```
data baseball; set sashelp.baseball;
proc contents varnum data=baseball;
   ods select position;
run;
```

Variables in Creation Order

#	Variable	Type	Len	Label
1	Name	Char	18	Player's Name
2	Team	Char	14	Team at the End of 1986
3	nAtBat	Num	8	Times at Bat in 1986
4	nHits	Num	8	Hits in 1986
5	nHome	Num	8	Home Runs in 1986
6	nRuns	Num	8	Runs in 1986
7	nRBI	Num	8	RBIs in 1986
8	nBB	Num	8	Walks in 1986
9	YrMajor	Num	8	Years in the Major Leagues
10	CrAtBat	Num	8	Career Times at Bat
11	CrHits	Num	8	Career Hits
12	CrHome	Num	8	Career Home Runs
13	CrRuns	Num	8	Career Runs
14	CrRbi	Num	8	Career RBIs
15	CrBB	Num	8	Career Walks
16	League	Char	8	League at the End of 1986
17	Division	Char	8	Division at the End of 1986
18	Position	Char	8	Position(s) in 1986
19	nOuts	Num	8	Put Outs in 1986
20	nAssts	Num	8	Assists in 1986
21	nError	Num	8	Errors in 1986
22	Salary	Num	8	1987 Salary in $ Thousands
23	Div	Char	16	League and Division
24	logSalary	Num	8	Log Salary

```
/* lasso */
proc glmselect data=baseball plots=(criterion ase);
 class league division;
 model logSalary = nAtBat nHits nHome nRuns nRBI nBB
                   yrMajor crAtBat crHits crHome crRuns crRbi
                   crBB league division nOuts nAssts nError
        / selection=lasso(adaptive choose=sbc stop=none);
 output out=out1 p=predlasso;
run;
```

Data Set	WORK.BASEBALL
Dependent Variable	logSalary
Selection Method	Adaptive LASSO
Stop Criterion	None
Choose Criterion	SBC
Effect Hierarchy Enforced	None

Number of Observations Read	322
Number of Observations Used	263

Class Level Information

Class	*Levels*	*Values*
League	2	American National
Division	2	East West

LASSO Selection Summary

Step	*Effect Entered*	*Effect Removed*	*Number Effects In*	*SBC*
		** Optimal Value of Criterion*		
0	*Intercept*		1	-57.2041
1	*CrRuns*		2	-70.8348
2	*nHits*		3	-226.0696
3	*CrHits*		4	-238.6648
4	*YrMajor*		5	-248.4971
5	*nBB*		6	-260.5682*
6	*Division_East*		7	-257.7020
7	*nOuts*		8	-254.3352
8		*CrRuns*	7	-260.1040
9	*nError*		8	-254.9990
10	*CrBB*		9	-249.9243
11	*CrRuns*		10	-245.7008
12	*nAtBat*		11	-241.6564
13	*nHome*		12	-236.3245
14	*League_American*		13	-238.1068
15	*CrAtBat*		14	-234.0015
16		*CrHits*	13	-241.0870
17	*nAssts*		14	-235.9894
18	*CrHits*		15	-230.5456
19	*nRuns*		16	-225.5197
20	*nRBI*		17	-220.3634
21	*CrRbi*		18	-214.7952
22	*CrHome*		19	-209.2505

Selection stopped because all candidate effects for entry are linearly dependent on effects in the model.

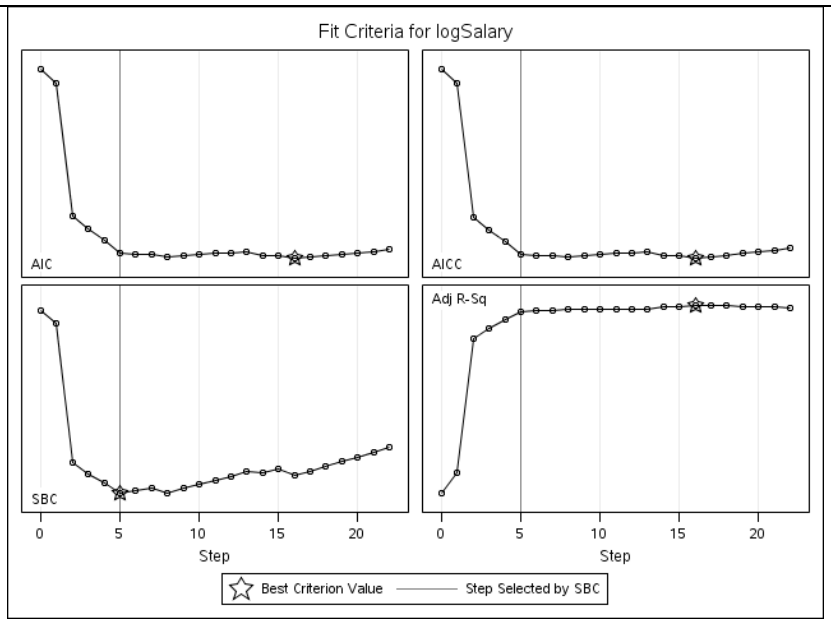

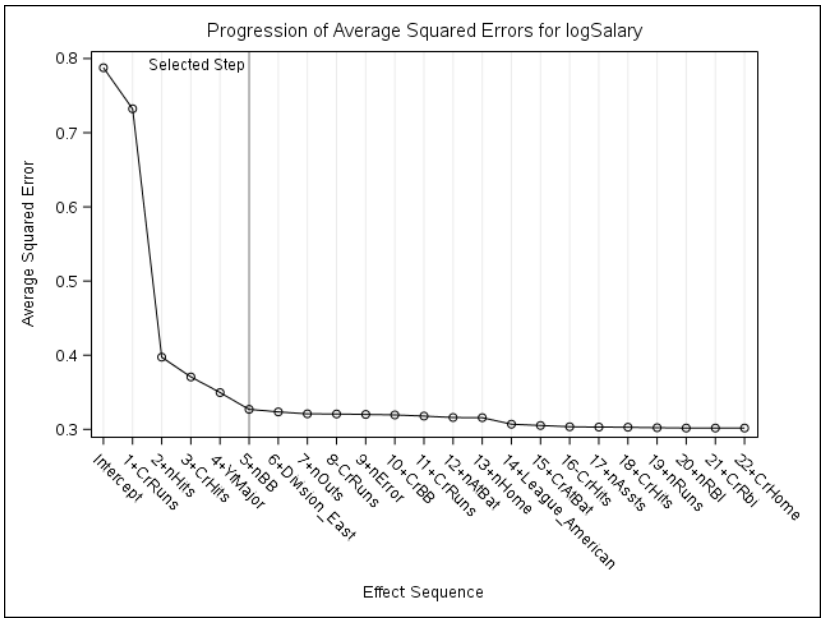

Selected Model

The selected model, based on SBC, is the model at Step 5.

Root MSE	0.57845
Dependent Mean	5.92722
R-Square	0.5849
Adj R-Sq	0.5768
AIC	-17.00115
AICC	-16.56194
SBC	-260.56823

Parameter Estimates

Parameter	*DF*	*Estimate*
Intercept	1	4.229778
nHits	1	0.007194
nBB	1	0.005629
YrMajor	1	0.062808
CrHits	1	0.000222
CrRuns	1	0.000136

```
/* elastic net */
proc glmselect data=out1 plots=(criterion ase) seed=12;
 class league division;
 model logSalary = nAtBat nHits nHome nRuns nRBI nBB
                   yrMajor crAtBat crHits crHome crRuns crRbi
                   crBB league division nOuts nAssts nError
      / selection=elasticnet(stop=none choose=cv)
        cvmethod=random(20);
 output out=out2 p=predelasticnet;
run;
```

Data Set	WORK.OUT1
Dependent Variable	logSalary
Selection Method	ELASTICNET
Stop Criterion	None
Choose Criterion	Cross Validation
Cross Validation Method	Random
Cross Validation Fold	20
Effect Hierarchy Enforced	None
Random Number Seed	12

Number of Observations Read	322
Number of Observations Used	263

Class Level Information

Class	*Levels*	*Values*
League	2	American National
Division	2	East West

Elastic Net Selection Summary

Step	*Effect Entered*	*Effect Removed*	*Number Effects In*	*CV PRESS*
		** Optimal Value of Criterion*		
0	*Intercept*		1	209.2326
1	*CrRuns*		2	123.1776
2	*CrHits*		3	123.7433
3	*nHits*		4	97.6956
4	*nBB*		5	94.7216
5	*CrRbi*		6	98.1015
6	*YrMajor*		7	92.7082
7	*nRBI*		8	94.5500
8	*Division_East*		9	93.3921
9	*nOuts*		10	94.1530
10	*nError*		11	93.8913
11	*nHome*		12	94.2533
12	*League_American*		13	94.4968
13		*CrRbi*	12	90.7314
14		*CrRuns*	11	90.1957
15		*nRBI*	10	89.6571
16	*CrBB*		11	89.2733
17	*CrRuns*		12	89.4515
18	*nAtBat*		13	88.9017*
19	*CrAtBat*		14	89.2818
20	*nAssts*		15	89.7926
21	*CrHome*		16	91.8598
22	*nRBI*		17	92.6309
23	*nRuns*		18	93.1973
24	*CrRbi*		19	94.5881

** Optimal Value of Criterion*

Selection stopped because all candidate effects for entry are linearly dependent on effects in the model.

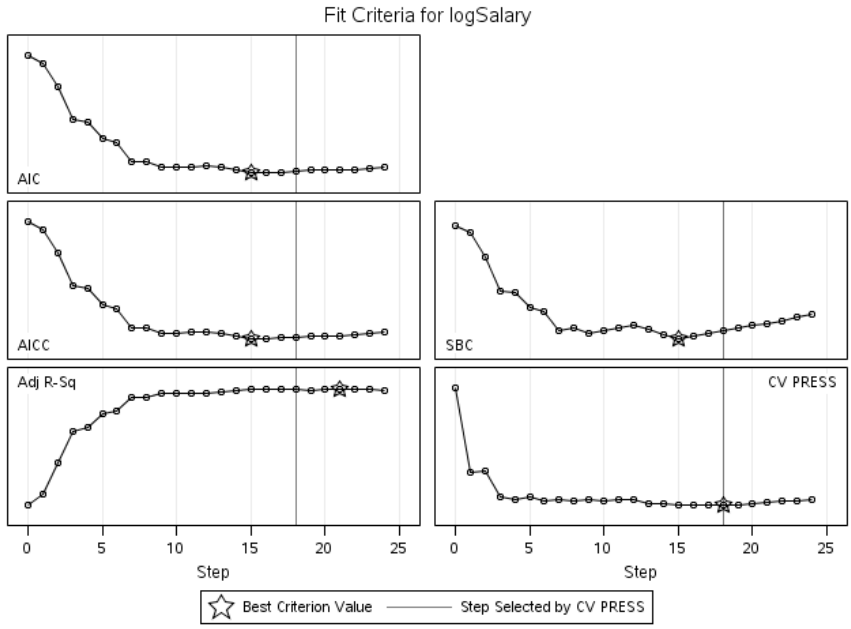

Selected Model
The selected model, based on Cross Validation, is the model at Step 18.

Root MSE	0.56923
Dependent Mean	5.92722
R-Square	0.6090
Adj R-Sq	0.5902
AIC	-18.72037
AICC	-17.02682
SBC	-237.28237
CV PRESS	88.90168

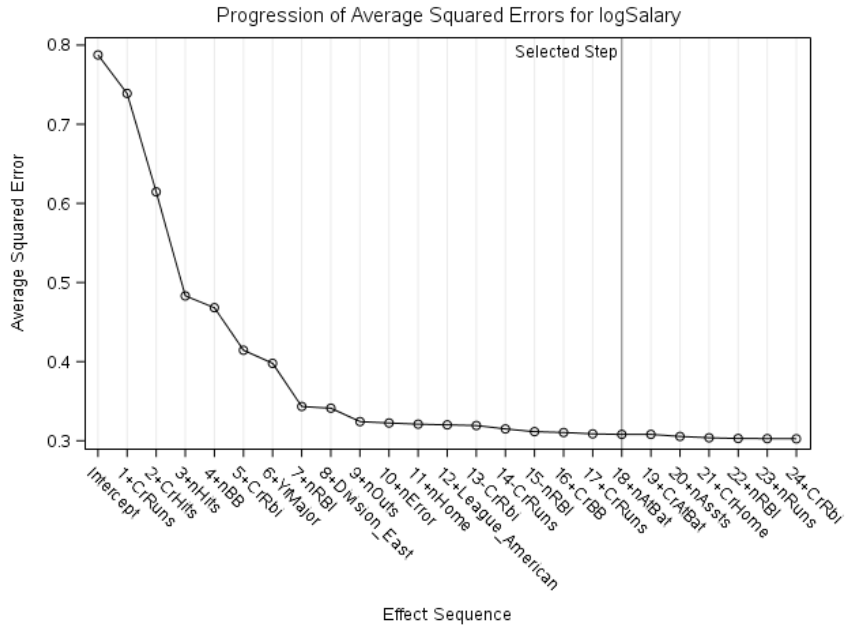

Parameter Estimates

Parameter	*DF*	*Estimate*
Intercept	1	4.195962
nAtBat	1	-0.000112
nHits	1	0.006807
nHome	1	0.003545
nBB	1	0.007082
YrMajor	1	0.070194
CrHits	1	0.000247
CrRuns	1	0.000212
CrBB	1	-0.000348
League_American	1	-0.092575
Division_East	1	0.144062
nOuts	1	0.000192
nError	1	-0.007767

```
proc sgscatter data=out2;
  matrix logSalary predlasso predelasticnet /
          markerattrs=(symbol=circlefilled size=6pt);
run;
```